

# COL764: Assignment 3

## IIT Delhi

Archisman Biswas, 2021MT10254

26 October 2024

## Contents

<b>1 Pointwise Learning to Rank (LETOR)</b>	<b>2</b>
1.1 Support Vector Regressor (SVR) . . . . .	2
1.2 Gradient Boosted Decision Tree (GBDT) . . . . .	2
1.3 Multi-Layer Perceptron (MLP) . . . . .	2
<b>2 Hyperparameter Tuning</b>	<b>3</b>
2.1 Support Vector Regressor . . . . .	3
2.1.1 TD2003 Dataset . . . . .	3
2.1.2 TD2004 Dataset . . . . .	5
2.2 Gradient Boosted Decision Tree . . . . .	6
2.2.1 TD2003 Dataset . . . . .	6
2.2.2 TD2004 Dataset . . . . .	8
2.3 Multi-Layer Perceptron . . . . .	10
2.3.1 TD2003 Dataset . . . . .	10
2.3.2 TD2004 Dataset . . . . .	12
<b>3 Feature Processing</b>	<b>14</b>
3.1 Support Vector Regressor . . . . .	14
3.2 Gradient Boosted Decision Tree . . . . .	15
3.3 Multi-Layer Perceptron . . . . .	16
<b>4 Results</b>	<b>18</b>
4.1 Support Vector Regressor . . . . .	18
4.2 Gradient Boosted Decision Trees . . . . .	19
4.3 Multi-Layer Perceptron . . . . .	20
<b>5 Bonus Task</b>	<b>21</b>
5.1 Hyperparameters . . . . .	21
5.2 Results . . . . .	21

# 1 Pointwise Learning to Rank (LETOR)

In pointwise learning to rank, each document-query pair is treated as an individual regression problem where a relevance score is predicted. The models minimize a loss function that measures the difference between the predicted and true relevance scores.

## 1.1 Support Vector Regressor (SVR)

The Support Vector Regressor solves the following optimization problem:

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n L_\epsilon(y_i, f(\mathbf{x}_i)),$$

where  $f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + b$  is the regression function,  $L_\epsilon$  is the  $\epsilon$ -insensitive loss function given by

$$L_\epsilon(y_i, f(\mathbf{x}_i)) = \max(0, |y_i - f(\mathbf{x}_i)| - \epsilon),$$

and  $C$  controls the trade-off between the flatness of the model and the amount up to which deviations larger than  $\epsilon$  are tolerated.

Library used for implementation: **SVR** from `sklearn.svm`

## 1.2 Gradient Boosted Decision Tree (GBDT)

Gradient Boosted Decision Trees optimize an ensemble of trees by minimizing a differentiable loss function  $L(y, \hat{y})$  over the model's predictions:

$$\hat{y} = \sum_{m=1}^M \gamma_m h_m(\mathbf{x}),$$

where  $M$  is the number of trees,  $h_m$  are the individual decision trees, and  $\gamma_m$  are the weights found by line search to minimize

$$L(y_i, \hat{y}_i^{(m-1)} + \gamma_m h_m(\mathbf{x}_i)),$$

where  $\hat{y}_i^{(m-1)}$  is the prediction from the first  $m - 1$  trees.

Library used for implementation: **GradientBoostingRegressor** from `sklearn.ensemble`

## 1.3 Multi-Layer Perceptron (MLP)

The Multi-Layer Perceptron is a feedforward neural network that minimizes a loss function  $L(y, \hat{y})$ , where the predicted relevance score is given by:

$$\hat{y} = \sigma \left( \mathbf{W}^{(L)} \sigma \left( \cdots \sigma \left( \mathbf{W}^{(1)} \mathbf{x} + \mathbf{b}^{(1)} \right) + \cdots + \mathbf{b}^{(L)} \right) \right),$$

with  $L$  layers, weight matrices  $\mathbf{W}^{(l)}$ , bias vectors  $\mathbf{b}^{(l)}$ , and an activation function  $\sigma$ , typically ReLU  $\sigma(z) = \max(0, z)$ .

Library used for implementation: **MLPRegressor** from `sklearn.neural_network`

## 2 Hyperparameter Tuning

Initially, the hyperparameters are set to their default values. Each hyperparameter is then tuned by evaluating the model's performance across a range of possible values. For that a **sequential search** is done on the parameters, where one at a time is tuned while the others are constant (at their default values). In this part, there is not much optimization for runtime required, since the number of training examples is on the lower side.

### Evaluation metric

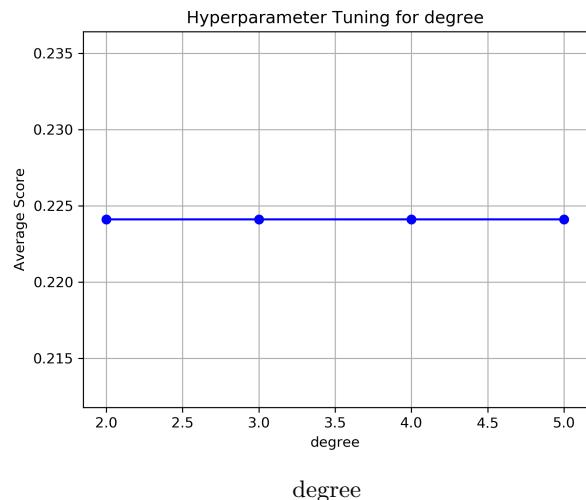
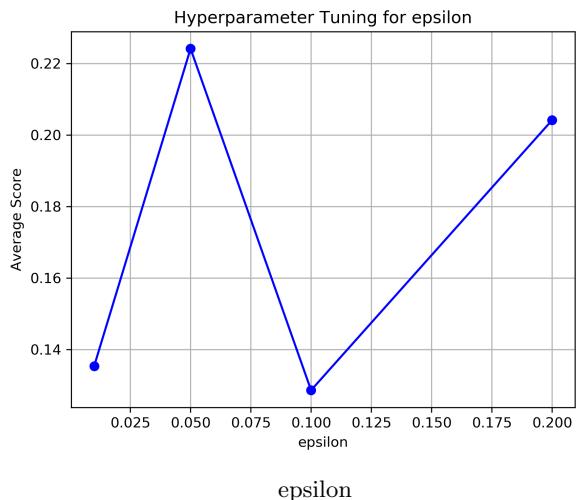
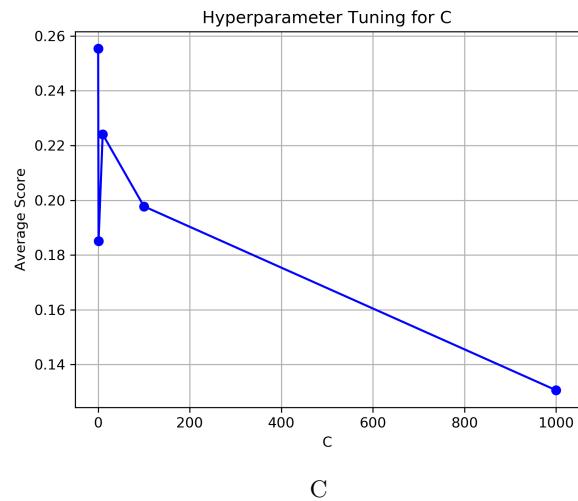
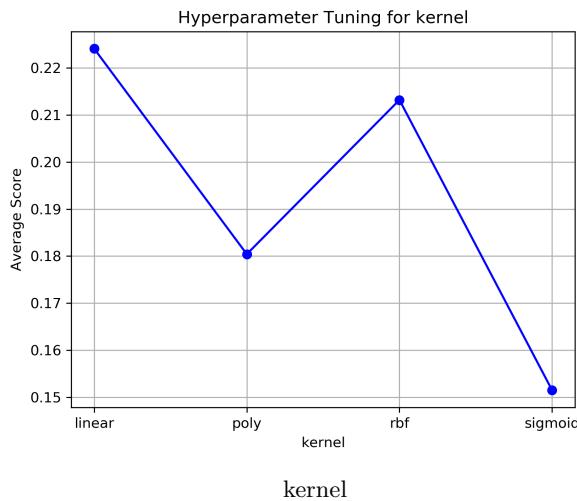
For each model, compute the average of  $\min_i \text{avg\_nDCG}[i]$  across the four folds, where  $\text{avg\_nDCG}[i]$  for a fold represents the nDCG[i] values averaged for a query.

### 2.1 Support Vector Regressor

kernel	C	epsilon	degree	max_iter	shrinking	gamma
'rbf'	1.0	0.1	3	-1	True	'auto'

Table 1: Hyperparameters to be tuned (with default values)

#### 2.1.1 TD2003 Dataset



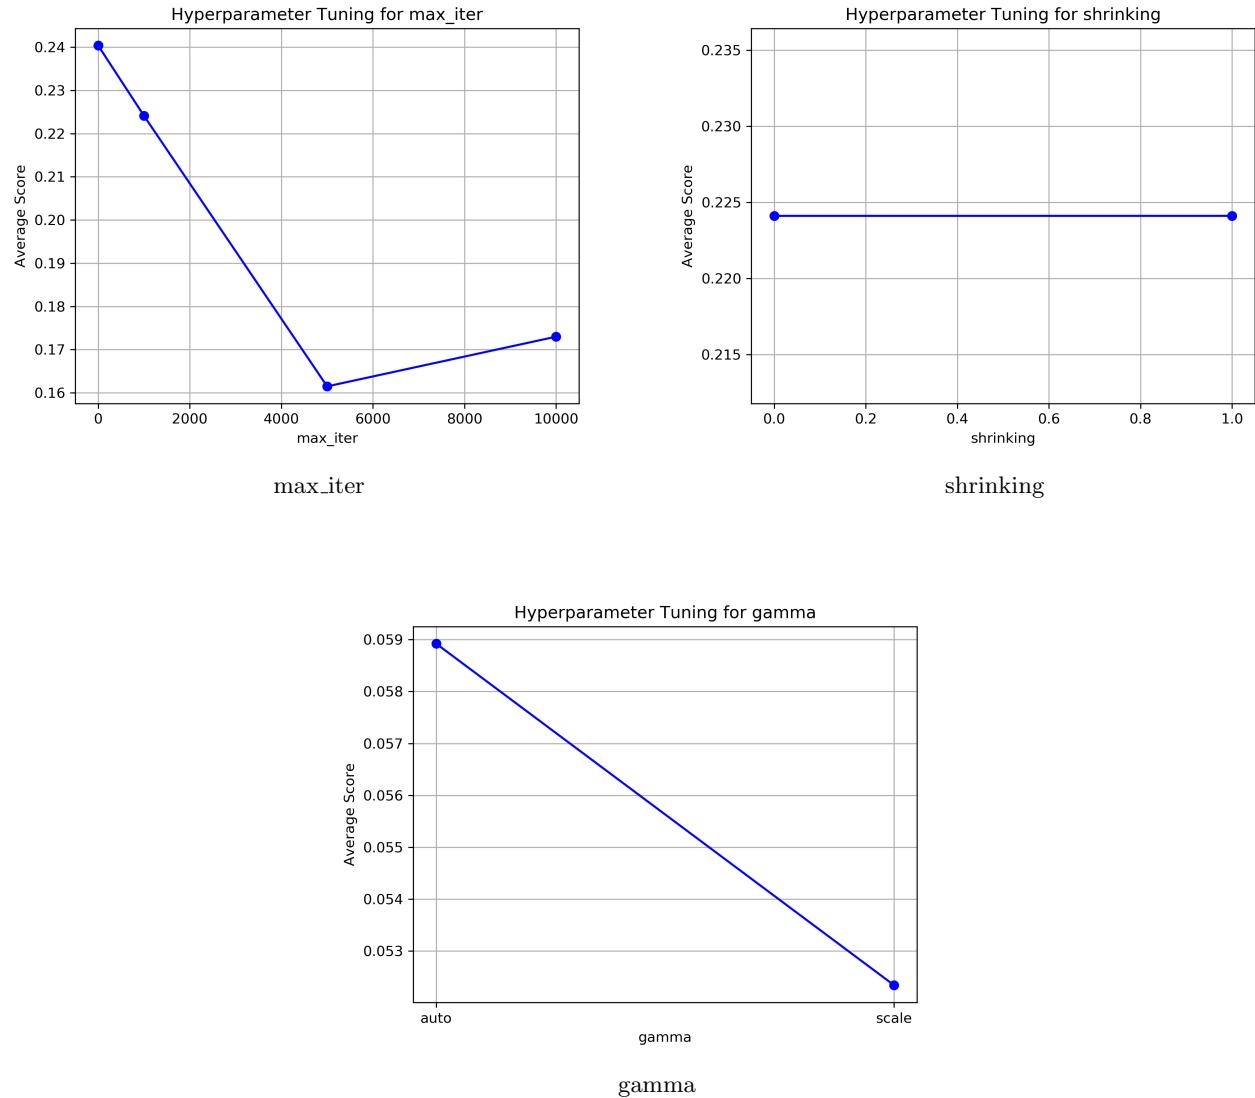


Figure 4: Results on different hyperparameter values

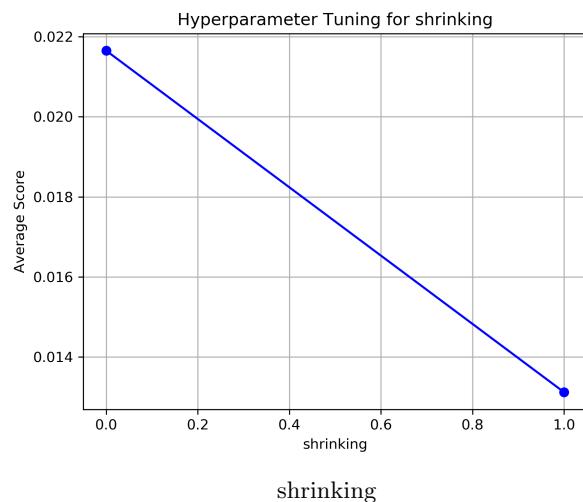
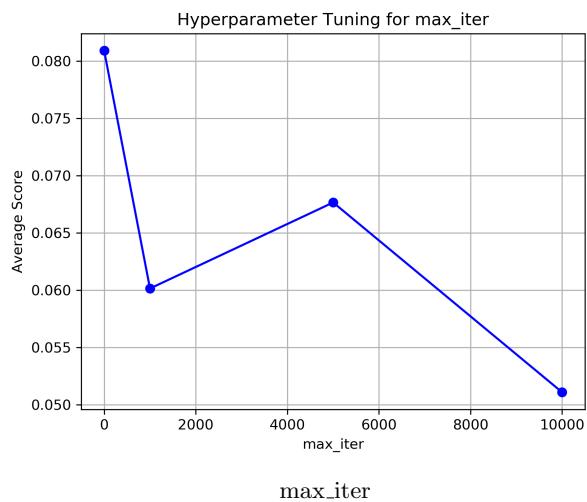
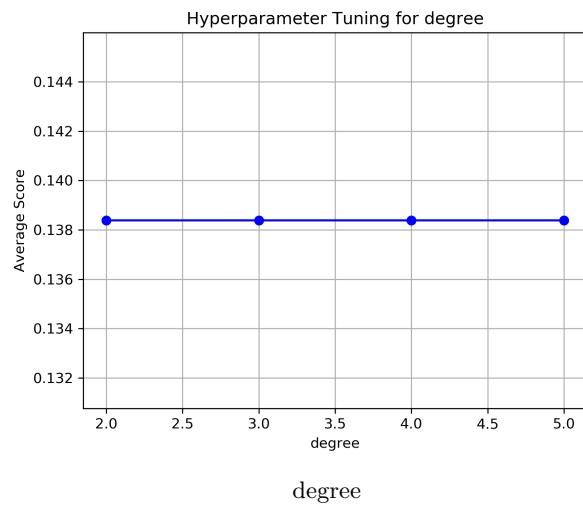
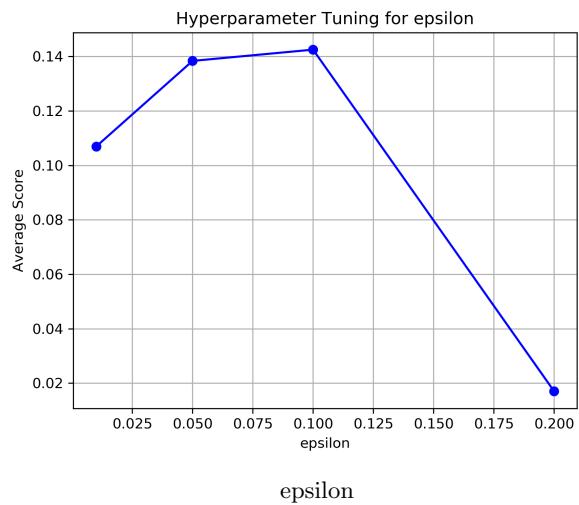
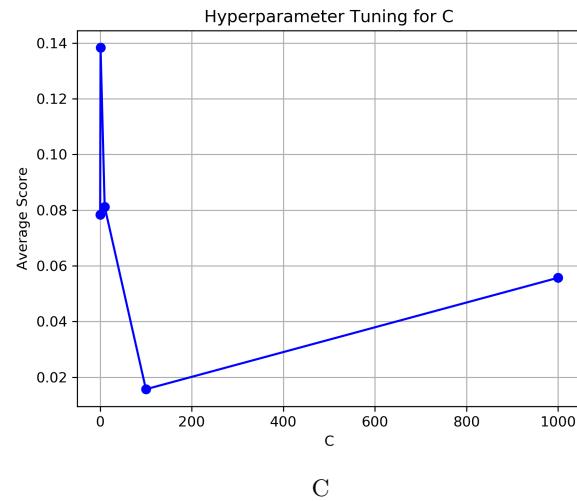
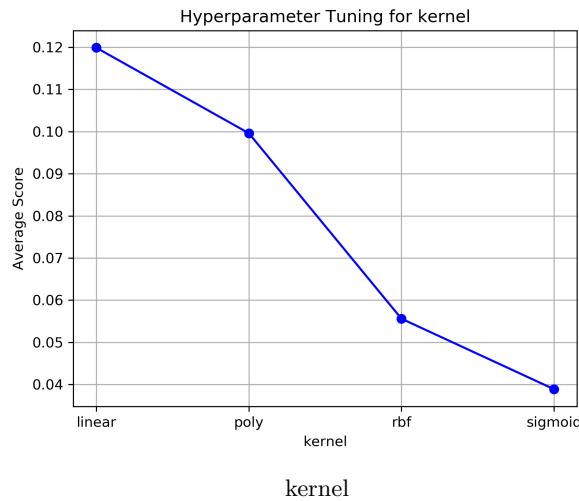
Therefore, based on the above results we choose the following values:

kernel	C	epsilon	degree	max_iter	shrinking	gamma
'Linear'	10.0	0.05	3	-1	True	'auto'

Table 2: SVR hyperparameters for TD2003

Note: C=10.0 has been chosen after taking into account the consistency in the pattern showed.

### 2.1.2 TD2004 Dataset



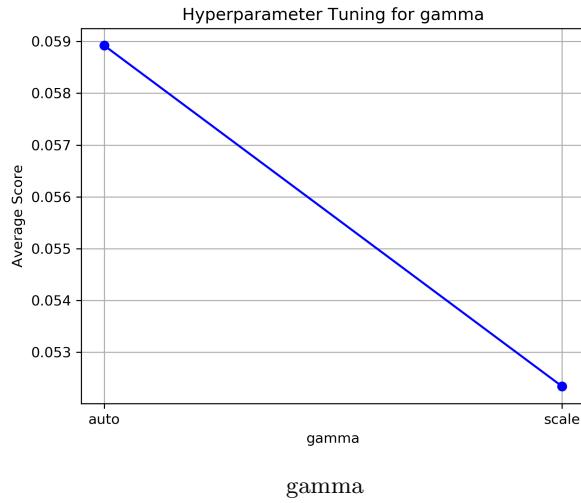


Figure 8: Results on different hyperparameter values

Therefore, based on the above results we choose the following values:

kernel	C	epsilon	degree	max_iter	shrinking	gamma
'Linear'	10.0	0.1	3	-1	True	'auto'

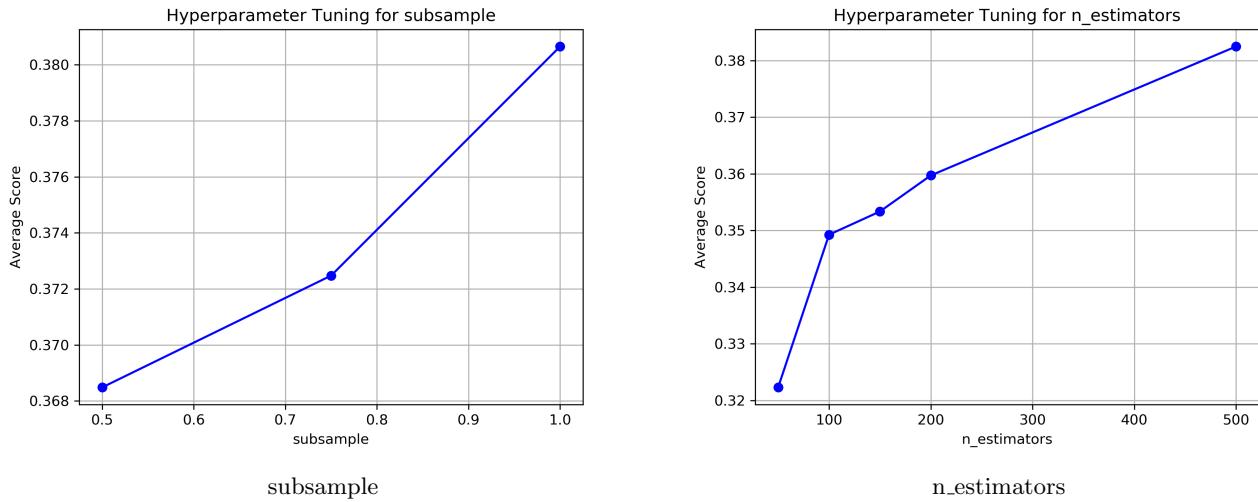
Table 3: SVR hyperparameters for TD2004

## 2.2 Gradient Boosted Decision Tree

subsample	n_estimators	learning_rate	max_depth	criterion <sup>1</sup>	random_state	max_features	alpha
1.0	100	0.1	3	'friedman_mse'	None	'auto'	0.0

Table 4: Hyperparameters to be tuned (with default values)

### 2.2.1 TD2003 Dataset



<sup>1</sup>Fixed

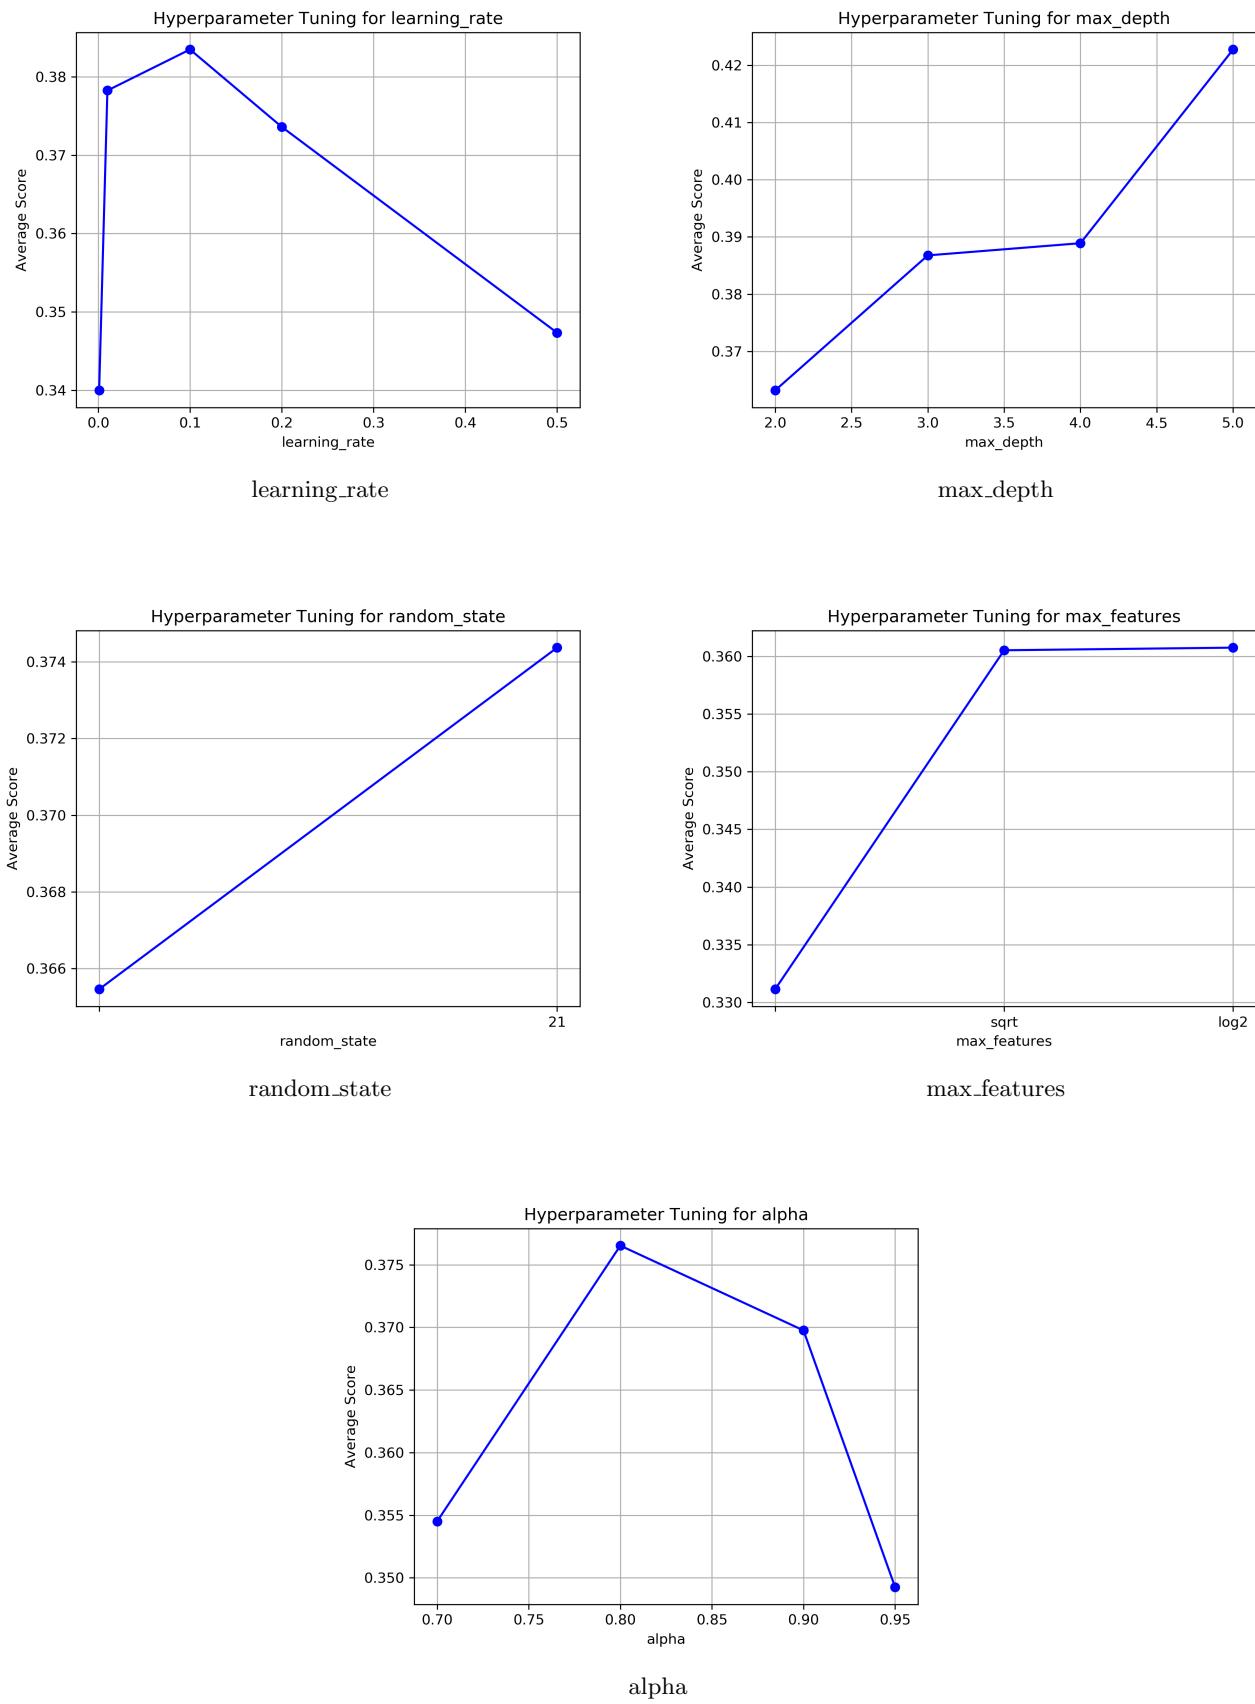


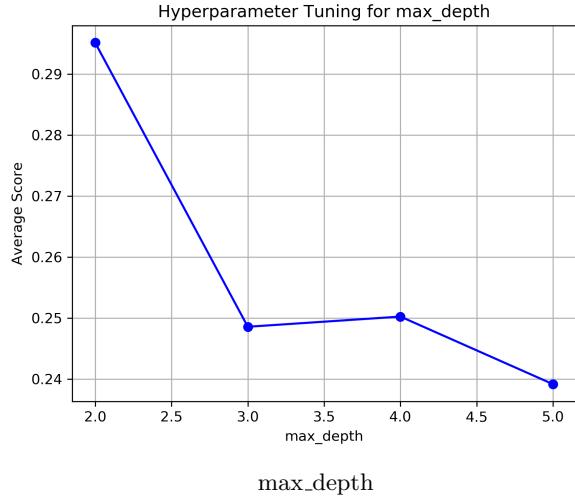
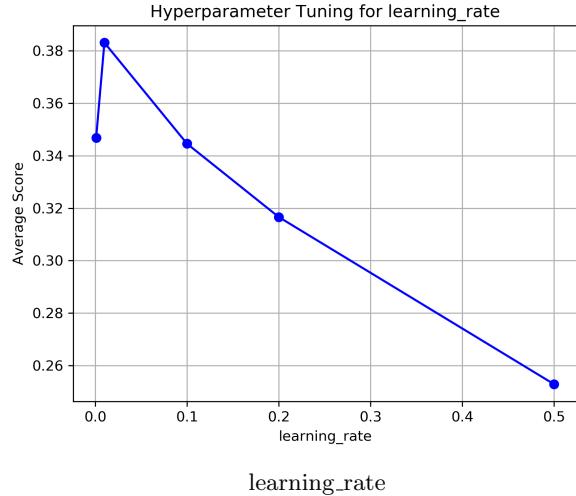
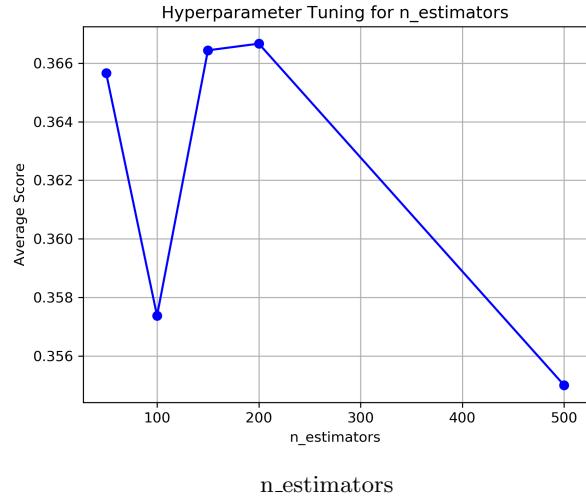
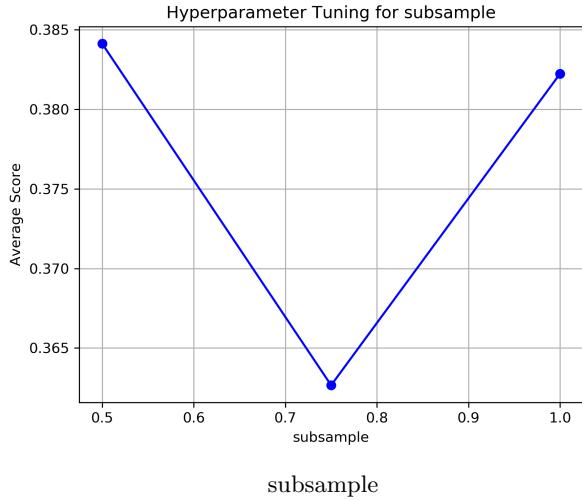
Figure 12: Results on different hyperparameter values

Therefore, based on the above results we choose the following values:

subsample	n_estimators	learning_rate	max_depth	criterion	random_state	max_features	alpha
1.0	500	0.1	5	'friedman_mse'	21	'sqrt'	0.8

Table 5: GBDT hyperparameters for TD2003

### 2.2.2 TD2004 Dataset



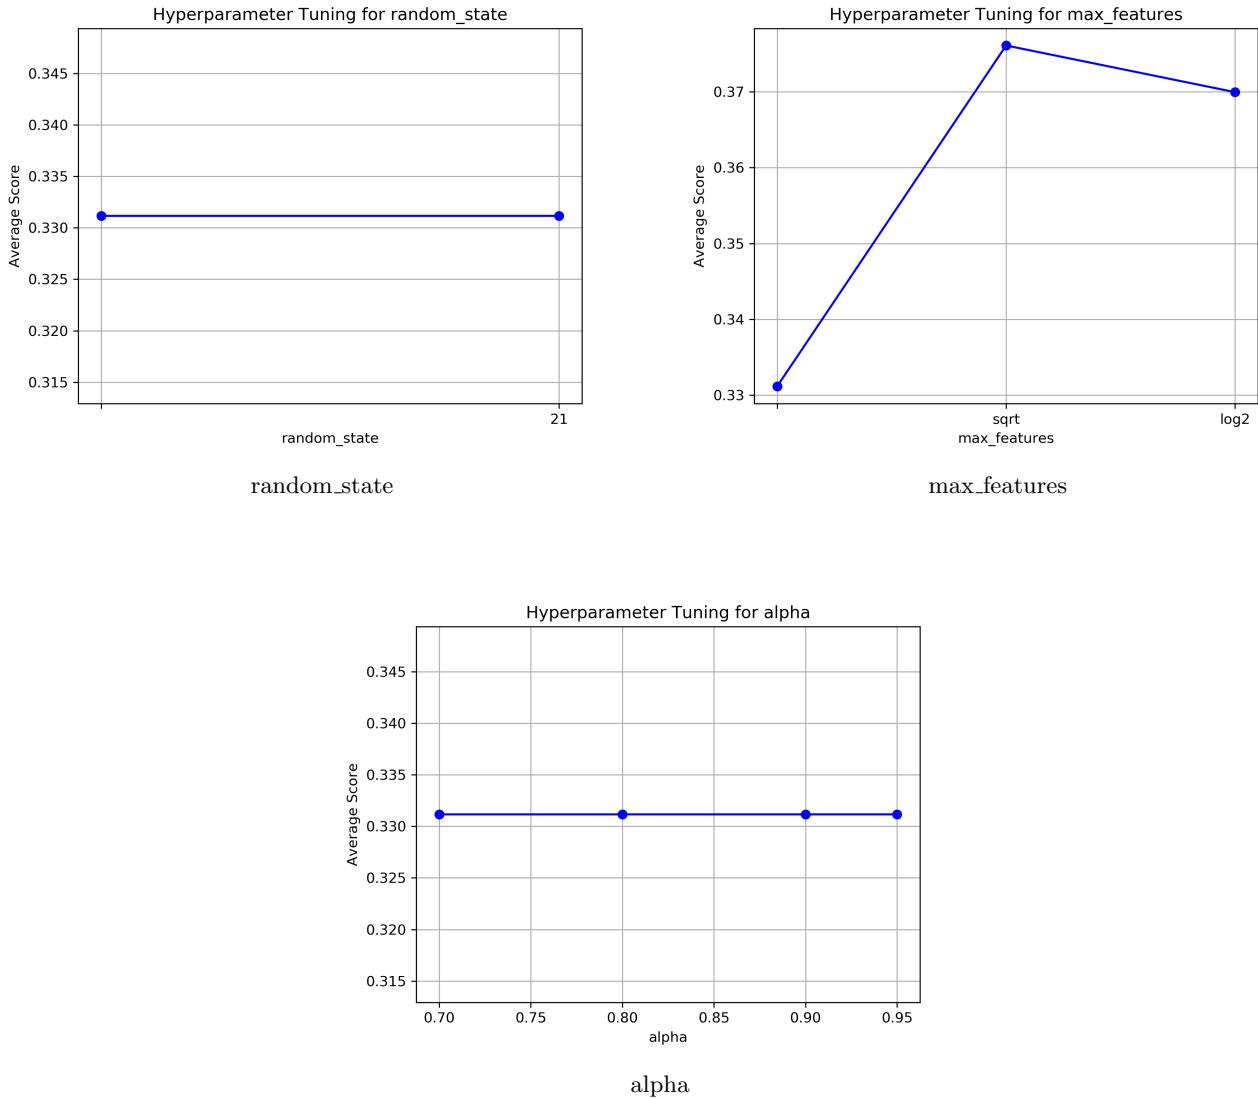


Figure 16: Results on different hyperparameter values

Therefore, based on the above results we choose the following values:

subsample	n_estimators	learning_rate	max_depth	criterion	random_state	max_features	alpha
1.0	200	0.01	2	'friedman_mse'	21	'sqrt'	0.8

Table 6: GBDT hyperparameters for TD2004

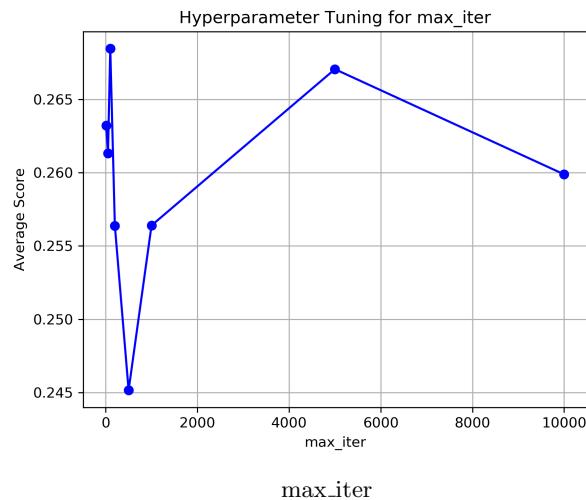
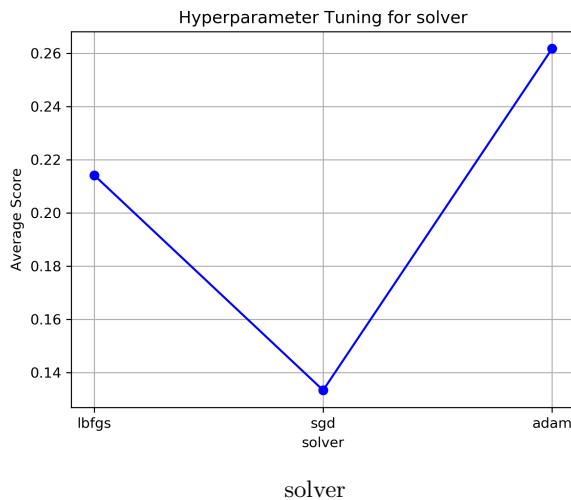
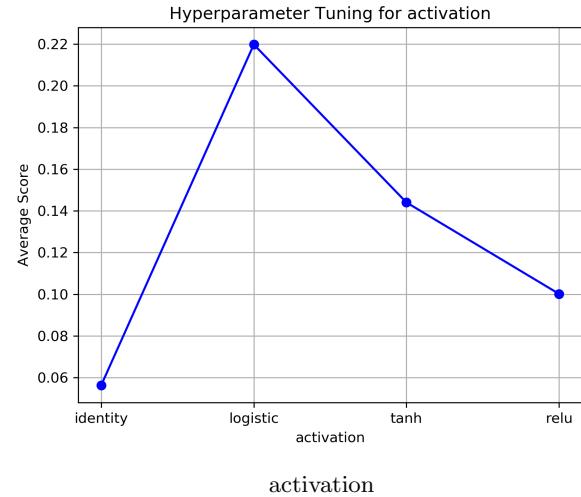
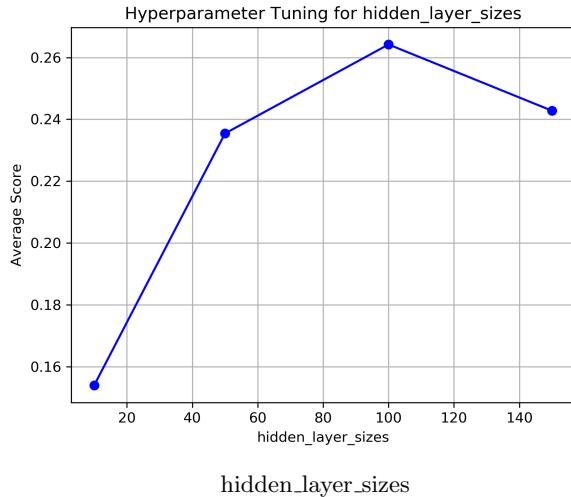
Note that ties are resolved based on results shown in **TD2003** dataset.

## 2.3 Multi-Layer Perceptron

hidden_layer_sizes	activation	solver	max_iter	alpha	learning_rate	learning_rate_init	power_t	random_state
(100,)	'relu'	'adam'	200	0.0001	'constant'	0.001	0.5	None

Table 7: Hyperparameters to be tuned (with default values)

### 2.3.1 TD2003 Dataset



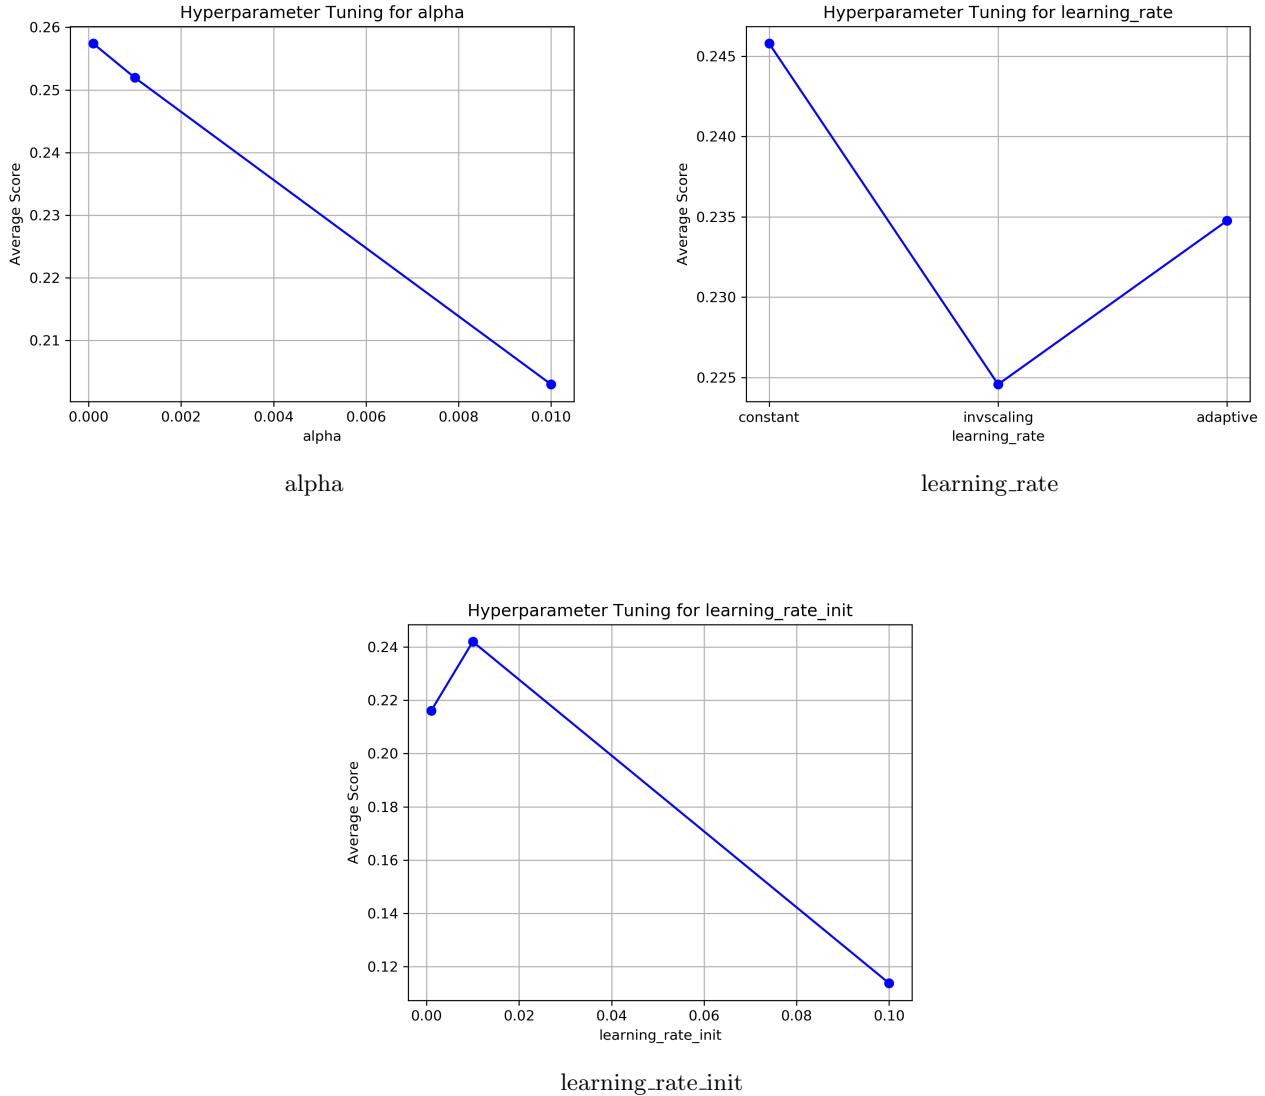


Figure 20: Results on different hyperparameter values

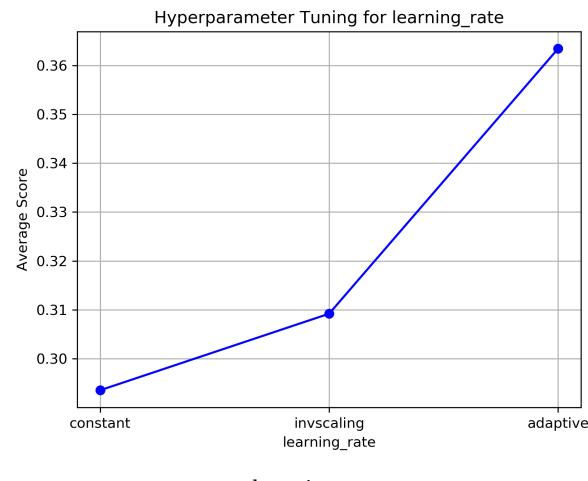
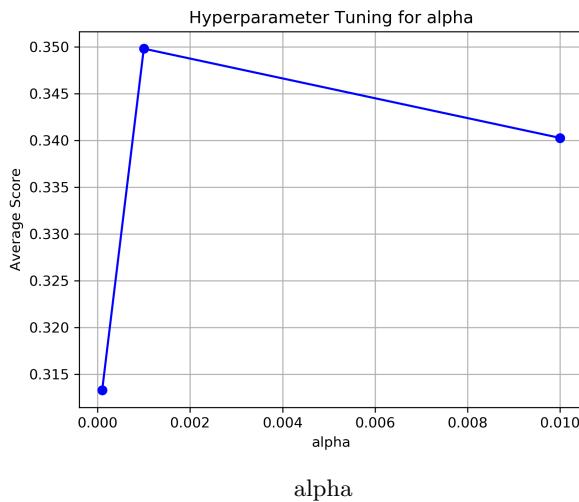
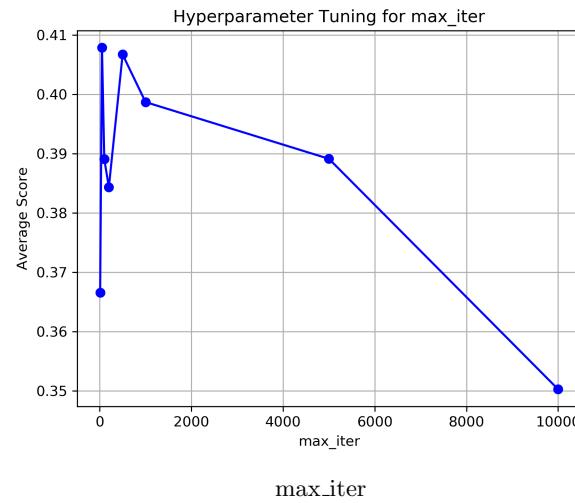
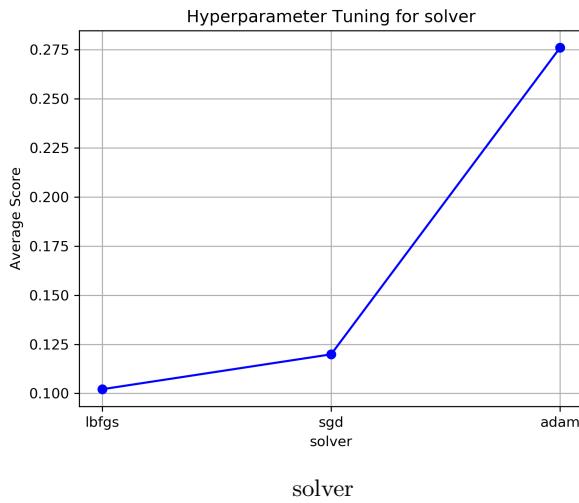
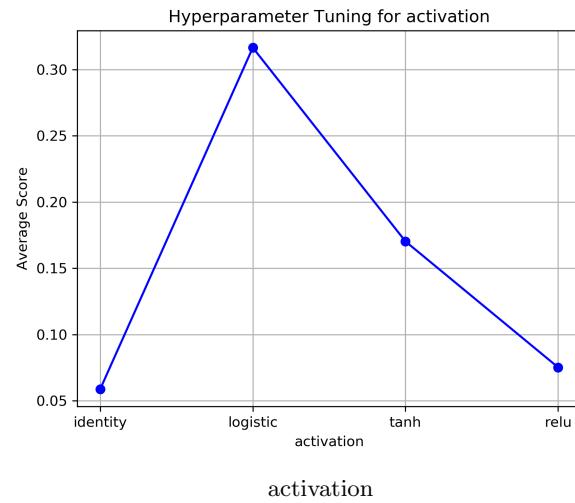
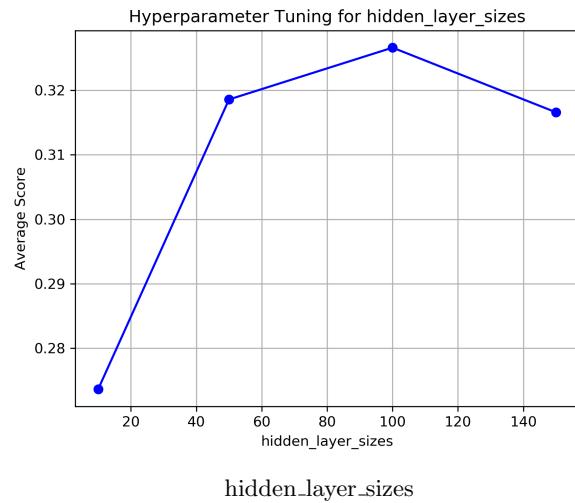
Therefore, based on the above results we choose the following values:

hidden_layer_sizes	activation	solver	max_iter	alpha	learning_rate	learning_rate_init	power_t	random_state
(100,)	'logistic'	'adam'	5000	0.0001	'adaptive'	0.01	0.5	60

Table 8: MLP hyperparameters for TD2003

Note that `max_iter`=200 has been chosen according to the pattern observed (in here as well in TD2004) and computation limits. `random_state` has been chosen to be some integer at random so that the quoted results are reproducible. `power_t`, the exponent for inverse scaling learning rate is fixed since it is only used when `solver='sgd'`. `learning_rate` has been set to **adaptive**, since it improves the runtime and does not differ too much to the maximum case. In that case, it will also be consistent with the results shown by TD2004.

### 2.3.2 TD2004 Dataset



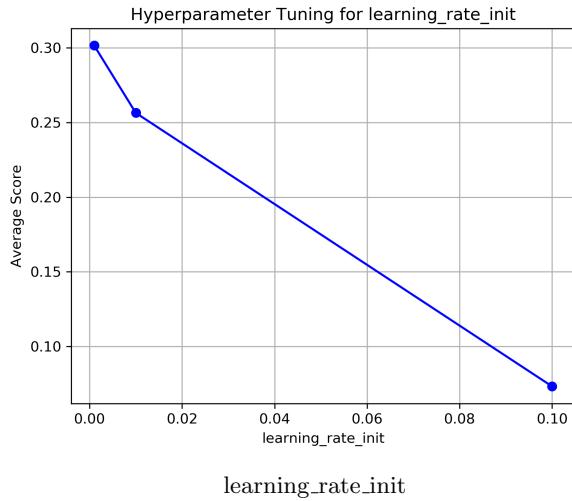


Figure 24: Results on different hyperparameter values

Therefore, based on the above results we choose the following values:

hidden_layer_sizes	activation	solver	max_iter	alpha	learning_rate	learning_rate_init	power_t	random_state
(100,)	'logistic'	'adam'	5000	0.001	'adaptive'	0.001	0.5	0

Table 9: MLP hyperparameters for TD2004

Note that `max_iter=5000` has been chosen to out of consistency with the previous part, and satisfactory performance in this part too. The `random_state` choice is again to ensure reproducibility of the results. `power_t`, the exponent for inverse scaling learning rate is fixed since it is only used when `solver='sgd'`.

### 3 Feature Processing

In this part, tried out feature scaling, feature engineering, dimensionality reduction etc.

For feature scaling, min-max and standard scaling has been tried out, out of which the former fared better.

`MinMaxScaler` has been used for implementation.

Feature engineering and dimensionality reduction details are as follows:

$$\begin{aligned}f_1 &= \log(\text{BM25 score}) \\f_2 &= \text{BM25 score of the anchor} \times \text{BM25 score of the title} \\f_3 &= \text{HITS authority} \times \text{HITS hub} \\f_4 &= \text{tfidf body} \times \text{tfidf anchor}\end{aligned}$$

<b>0</b>	Original
<b>1</b>	Add feature $f_1$
<b>2</b>	Add features $f_1, f_2, f_3, f_4$
<b>3</b>	200 Add features $f_2, f_3, f_4$
<b>4</b>	0.001 Add feature $f_2$
<b>5</b>	Apply <code>PCA</code> (n_components=30)
<b>6</b>	Apply <code>t-SNE</code> (n_components=2)
<b>7</b>	Apply <code>truncatedSVD</code> (n_components=30)
<b>8</b>	Apply <code>FastICA</code> (n_components=30)

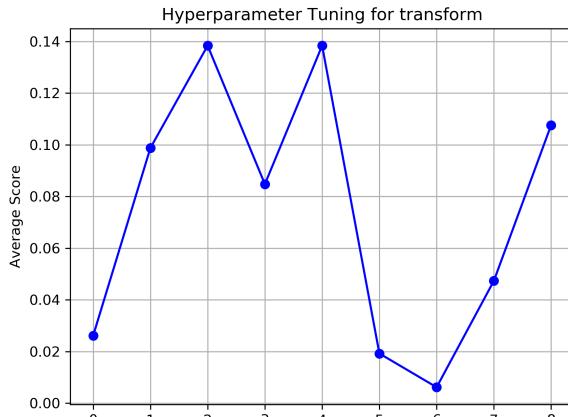
Table 10: Codes used for feature engineering and dimensionality reduction

**1, 2, 3, 4** is feature engineering, while **5, 6, 7, 8** is dimensionality reduction. Both cannot be done together since the original features lose their meaning after dimensionality reduction. Hence, all the 9 cases could be compared on a single graph.

#### Note

**6** is removed from the experiments due to very low evaluation score and long computation time

#### 3.1 Support Vector Regressor



TD2003

Figure 25: Experimentation results of SVR before min-max scaling

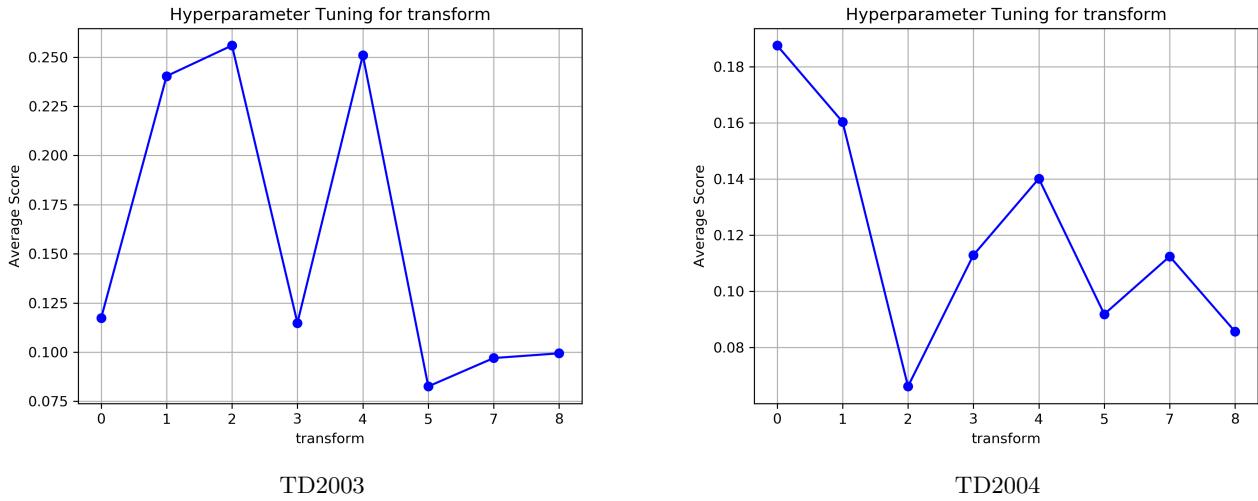


Figure 26: Experimentation results of SVR after min-max scaling

For scaling, we hence conclude that min-max scaling is much **better** for SVR. For feature engineering, after averaging out the results in both datasets, we notice that **1** (Adding feature  $f_1$ ) suits the best.

### 3.2 Gradient Boosted Decision Tree

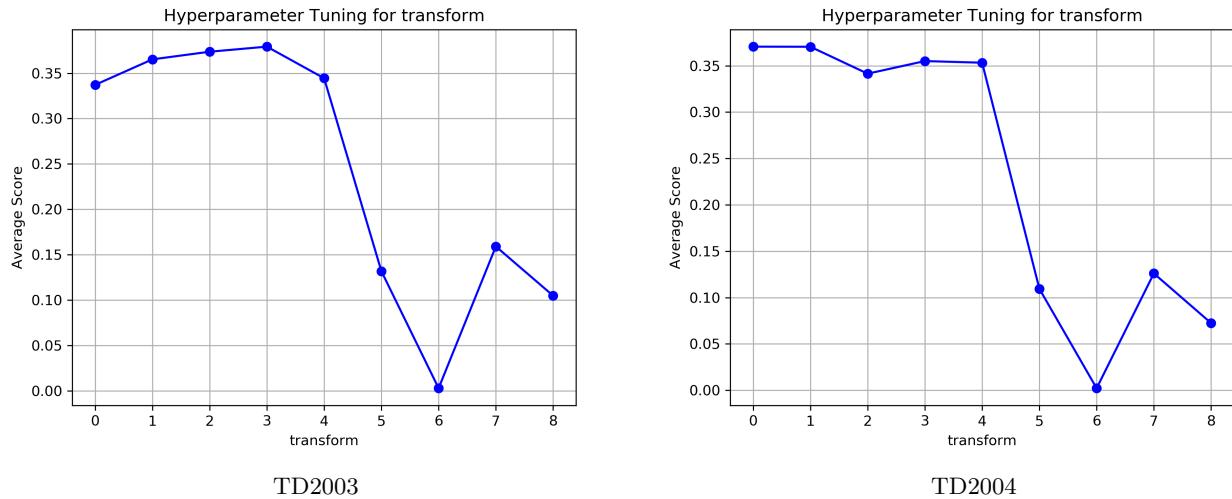
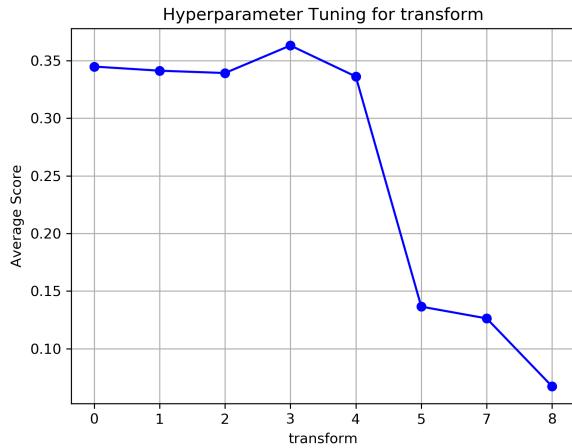


Figure 27: Experimentation results of MLP before min-max scaling

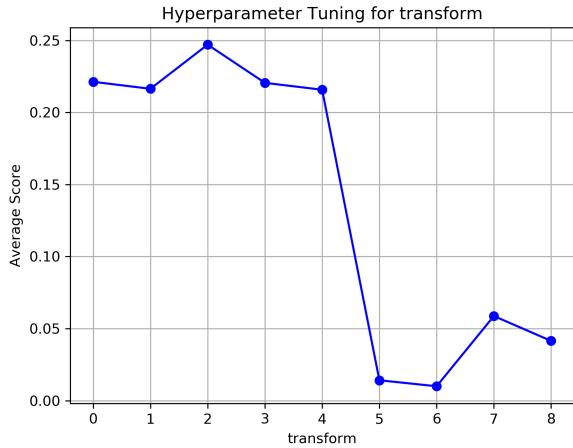


TD2003

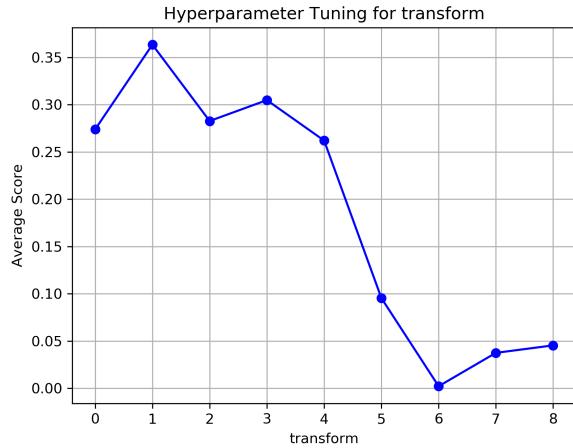
Figure 28: Experimentation results of GBDT after min-max scaling

Decision Trees are known to be prone to scaling and hence, as it can be seen min-max scaling **does not** work well. For feature engineering, after averaging out the results in both datasets, we notice that **1** (Adding feature  $f_1$ ) suits the best.

### 3.3 Multi-Layer Perceptron

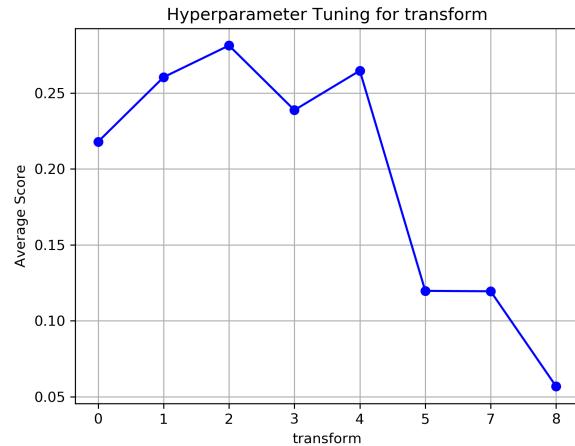


TD2003

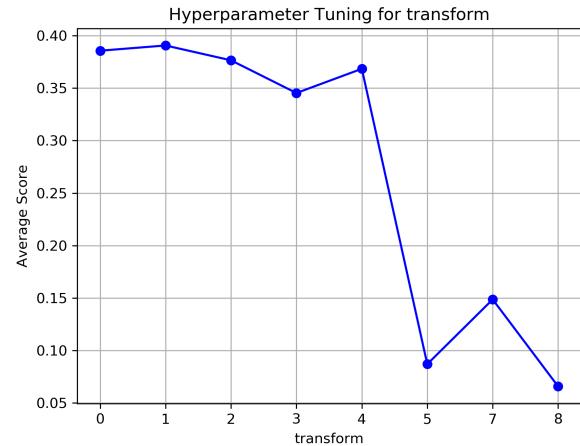


TD2004

Figure 29: Experimentation results of MLP before min-max scaling



TD2003



TD2004

Figure 30: Experimentation results of MLP after min-max scaling

For scaling, it is not satisfactorily clear if max-scaling is that good, so it is **not** done. For feature engineering, after averaging out the results in both datasets, we notice that **1** (Adding feature  $f_1$ ) suits the best.

## 4 Results

### Note

All following measures are given w.r.t programs run on Apple Silicon M2 chip

### 4.1 Support Vector Regressor

Fold 1			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.2555	0.2583	0.3399
<b>val</b>	0.4069	0.3778	0.3413
<b>test</b>	0.0339	0.0220	0.0352
Fold 2			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.2436	0.2197	0.2327
<b>val</b>	0.2276	0.1477	0.1213
<b>test</b>	0.0214	0.0139	0.0277
Fold 3			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.2698	0.2350	0.2328
<b>val</b>	0.3276	0.3275	0.3782
<b>test</b>	0.1069	0.0863	0.0822
Fold 4			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.2670	0.2212	0.2336
<b>val</b>	0.1946	0.2362	0.3491
<b>test</b>	0.2230	0.2144	0.2216

Table 11: nDCG results for TD2003

	Train Time (s)	Test Time (s)
<b>Fold 1</b>	8.7133	0.2224
<b>Fold 2</b>	20.8312	0.3307
<b>Fold 3</b>	16.5113	0.1682
<b>Fold 4</b>	19.0213	0.2418

Table 12: Train and test times for TD2003

Fold 1			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.1808	0.1876	0.2646
<b>val</b>	0.1244	0.1058	0.1622
<b>test</b>	0.0000	0.0000	0.0106
Fold 2			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.2408	0.2600	0.3795
<b>val</b>	0.0753	0.0797	0.1238
<b>test</b>	0.0000	0.0000	0.0197
Fold 3			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.2345	0.2702	0.3753
<b>val</b>	0.0498	0.0511	0.1058
<b>test</b>	0.0185	0.0120	0.0408
Fold 4			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.2196	0.2448	0.3425
<b>val</b>	0.2780	0.2676	0.3957
<b>test</b>	0.0000	0.0000	0.0200

Table 13: nDCG results for TD2004

	Train Time (s)	Test Time (s)
<b>Fold 1</b>	47.5260	0.3589
<b>Fold 2</b>	48.3123	0.3662
<b>Fold 3</b>	88.5335	0.3064
<b>Fold 4</b>	53.4189	0.2962

Table 14: Train and test times for TD2004

## 4.2 Gradient Boosted Decision Trees

Fold 1			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	1.0000	1.0000	0.9999
<b>val</b>	0.4213	0.3822	0.3681
<b>test</b>	0.2700	0.2467	0.3080
Fold 2			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.9994	0.9922	0.9953
<b>val</b>	0.3661	0.3266	0.3679
<b>test</b>	0.2369	0.2303	0.2987
Fold 3			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	1.0000	0.9910	0.9918
<b>val</b>	0.1863	0.1732	0.2610
<b>test</b>	0.2406	0.2426	0.3281
Fold 4			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.9877	0.9809	0.9823
<b>val</b>	0.1489	0.1917	0.3115
<b>test</b>	0.2552	0.2943	0.4093

Table 15: nDCG results for TD2003

	Train Time (s)	Test Time (s)
<b>Fold 1</b>	12.1395	0.0395
<b>Fold 2</b>	12.3275	0.0275
<b>Fold 3</b>	12.9135	0.0135
<b>Fold 4</b>	13.5633	0.0633

Table 16: Train and test times for TD2003

Fold 1			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.5046	0.5274	0.6226
<b>val</b>	0.3221	0.3483	0.4859
<b>test</b>	0.4531	0.4227	0.5038
Fold 2			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.5290	0.5381	0.6343
<b>val</b>	0.4208	0.4166	0.5057
<b>test</b>	0.2619	0.3278	0.4395
Fold 3			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.4444	0.4701	0.5601
<b>val</b>	0.2921	0.3489	0.4487
<b>test</b>	0.3894	0.4247	0.5299
Fold 4			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.4746	0.4685	0.5686
<b>val</b>	0.4667	0.4952	0.5965
<b>test</b>	0.3141	0.3388	0.4869

Table 17: nDCG results for TD2004

	Train Time (s)	Test Time (s)
<b>Fold 1</b>	3.2815	0.0815
<b>Fold 2</b>	3.2546	0.0546
<b>Fold 3</b>	3.6659	0.0659
<b>Fold 4</b>	3.5339	0.0339

Table 18: Train and test times for TD2004

### 4.3 Multi-Layer Perceptron

Fold 1			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.3231	0.3076	0.3969
<b>val</b>	0.3365	0.3204	0.3103
<b>test</b>	0.2935	0.2802	0.3540
Fold 2			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.4828	0.4240	0.4991
<b>val</b>	0.3008	0.2826	0.3960
<b>test</b>	0.1817	0.1799	0.2872
Fold 3			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.4699	0.4096	0.4636
<b>val</b>	0.1384	0.1243	0.2400
<b>test</b>	0.1612	0.2283	0.3189
Fold 4			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.4380	0.3575	0.4402
<b>val</b>	0.1447	0.2104	0.3228
<b>test</b>	0.2700	0.3007	0.3582

Table 19: nDCG results for TD2003

	Train Time (s)	Test Time (s)
<b>Fold 1</b>	0.4859	0.0138
<b>Fold 2</b>	0.4884	0.0153
<b>Fold 3</b>	0.5116	0.0131
<b>Fold 4</b>	0.5297	0.0163

Table 20: Train and test times for TD2003

Fold 1			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.4665	0.4909	0.5906
<b>val</b>	0.3099	0.2942	0.4352
<b>test</b>	0.3934	0.4369	0.4960
Fold 2			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.4285	0.4472	0.5596
<b>val</b>	0.4543	0.4559	0.5220
<b>test</b>	0.3270	0.3797	0.4775
Fold 3			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.3375	0.3720	0.4844
<b>val</b>	0.3701	0.3875	0.4712
<b>test</b>	0.3615	0.3855	0.5002
Fold 4			
	nDCG[5]	nDCG[10]	nDCG[100]
<b>train</b>	0.4086	0.4128	0.4969
<b>val</b>	0.3649	0.3827	0.4973
<b>test</b>	0.2129	0.2530	0.3633

Table 21: nDCG results for TD2004

	Train Time (s)	Test Time (s)
<b>Fold 1</b>	0.7266	0.0207
<b>Fold 2</b>	0.7213	0.0204
<b>Fold 3</b>	0.7468	0.0192
<b>Fold 4</b>	0.7324	0.0216

Table 22: Train and test times for TD2004

## 5 Bonus Task

In this part, a few adjustments have been done. First of which, only the **first 60** features for each query-document pair vector. This greatly helps in the runtime as well as the accuracy is also good. Some other methods like **PCA** was also tried after incorporating all features, but that did not really yield much satisfactory results. The following section details out the hyperparameters used and where they differ from the ones used in the previous parts.

### 5.1 Hyperparameters

In the **Support Vector Regressor** model, the dimensionality is reduced to **50** using **PCA**, resulting in an improved nDCG value. Additional optimizations, such as setting `max_iter=1000` to address runtime concerns, were also implemented. The final average nDCG value demonstrates satisfactory performance following these optimizations.

<code>kernel</code>	<code>C</code>	<code>epsilon</code>	<code>degree</code>	<code>max_iter</code>	<code>shrinking</code>	<code>gamma</code>
'Linear'	1.0	0.001	3	1000	True	'auto'

Table 23: SVR hyperparameters

In the **Gradient Boosted Decision Tree** model, due to the increased training size, `n_estimators` has been reduced to **100**, while `learning_rate` has been increased to **0.1**. These adjustments enhance runtime performance and yield a satisfactory average nDCG.

<code>subsample</code>	<code>n_estimators</code>	<code>learning_rate</code>	<code>max_depth</code>	<code>criterion</code>	<code>random_state</code>	<code>max_features</code>	<code>alpha</code>
1.0	200	0.1	2	'friedman_mse'	21	'sqrt'	0.8

Table 24: GBDT hyperparameters

In the **Multi-Layer Perceptron** model, no scaling is applied. The primary changes compared to the previous parameters are a reduction in the maximum number of iterations to **500** and an increase in the initial learning rate to **0.01**. These adjustments were made because the number of examples is significantly larger, which helps improve the runtime while still achieving satisfactory results, with an average nDCG greater than **0.7**.

<code>hidden_layer_sizes</code>	<code>activation</code>	<code>solver</code>	<code>max_iter</code>	<code>alpha</code>	<code>learning_rate</code>	<code>learning_rate_init</code>	<code>power_t</code>	<code>random_state</code>
(100,)	'logistic'	'adam'	500	0.01	'adaptive'	0.001	0.5	0

Table 25: MLP hyperparameters

### 5.2 Results

	<code>nDCG[5]</code>	<code>nDCG[10]</code>	<code>nDCG[100]</code>
<b>SVR</b>	0.5554	0.6269	0.7688
<b>GBDT</b>	0.6792	0.7250	0.8252
<b>MLP</b>	0.6821	0.7283	0.8281

Table 26: nDCG results on validation set

	<b>Train Time (s)</b>	<b>Validation Time (s)</b>
<b>SVR</b>	19.7673	3.0059
<b>GBDT</b>	86.0809	0.2542
<b>MLP</b>	57.9594	0.1773

Table 27: Time taken by the models

## List of Tables

1	Hyperparameters to be tuned (with default values) . . . . .	3
2	SVR hyperparameters for TD2003 . . . . .	4
3	SVR hyperparameters for TD2004 . . . . .	6
4	Hyperparameters to be tuned (with default values) . . . . .	6
5	GBDT hyperparameters for TD2003 . . . . .	8
6	GBDT hyperparameters for TD2004 . . . . .	9
7	Hyperparameters to be tuned (with default values) . . . . .	10
8	MLP hyperparameters for TD2003 . . . . .	11
9	MLP hyperparameters for TD2004 . . . . .	13
10	Codes used for feature engineering and dimensionality reduction . . . . .	14
11	nDCG results for TD2003 . . . . .	18
12	Train and test times for TD2003 . . . . .	18
13	nDCG results for TD2004 . . . . .	18
14	Train and test times for TD2004 . . . . .	18
15	nDCG results for TD2003 . . . . .	19
16	Train and test times for TD2003 . . . . .	19
17	nDCG results for TD2004 . . . . .	19
18	Train and test times for TD2004 . . . . .	19
19	nDCG results for TD2003 . . . . .	20
20	Train and test times for TD2003 . . . . .	20
21	nDCG results for TD2004 . . . . .	20
22	Train and test times for TD2004 . . . . .	20
23	SVR hyperparameters . . . . .	21
24	GBDT hyperparameters . . . . .	21
25	MLP hyperparameters . . . . .	21
26	nDCG results on validation set . . . . .	21
27	Time taken by the models . . . . .	21

## List of Figures

4	Results on different hyperparameter values . . . . .	4
8	Results on different hyperparameter values . . . . .	6
12	Results on different hyperparameter values . . . . .	7
16	Results on different hyperparameter values . . . . .	9
20	Results on different hyperparameter values . . . . .	11
24	Results on different hyperparameter values . . . . .	13
25	Experimentation results of SVR before min-max scaling . . . . .	14
26	Experimentation results of SVR after min-max scaling . . . . .	15
27	Experimentation results of MLP before min-max scaling . . . . .	15
28	Experimentation results of GBDT after min-max scaling . . . . .	16
29	Experimentation results of MLP before min-max scaling . . . . .	16
30	Experimentation results of MLP after min-max scaling . . . . .	17